# CS 643 Programming Assignment-2

mb2332@njit.edu
Manas Bhut

Github Link
https://github.com/Manas1227/cs643-853-pa2-mb2332/tree/master

Docker Link
https://hub.docker.com/repository/docker/manasbhut/cs643-pa2-aws-spark/general

## SECTION 1: AWS Cloud setup for running the training ML application - training.py

### Step-1: Create s3 bucket and upload required files

After creating new bucket, click on Add Files and upload both the datasets files to newly created S3 bucket

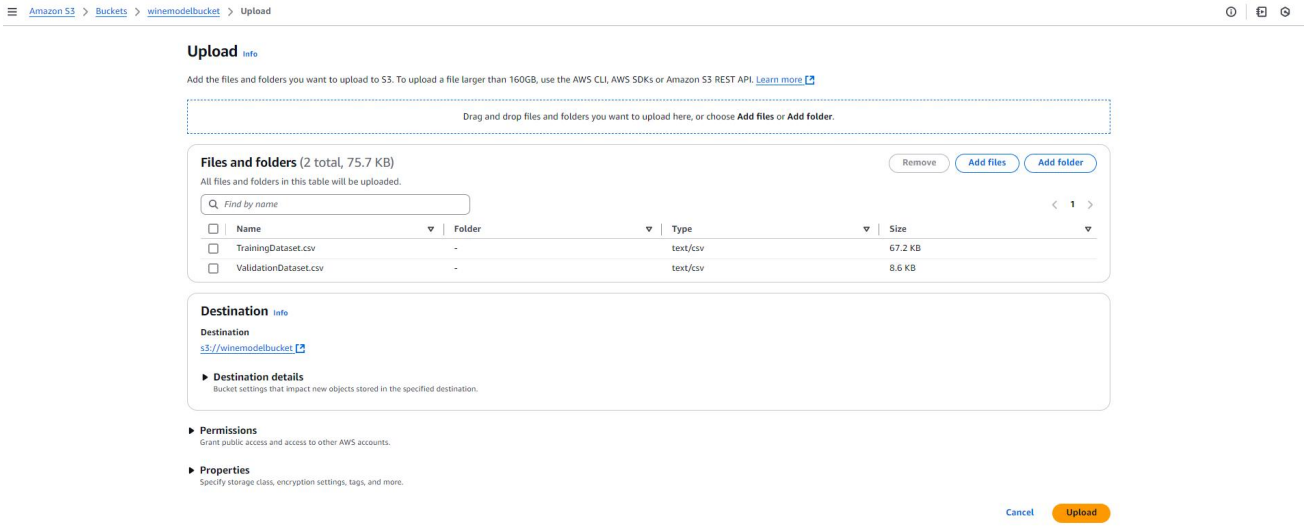

**Step-2 Create an EMR cluster**

Login to AWS console
Search for “EMR”
Click on “Create CLuster” option

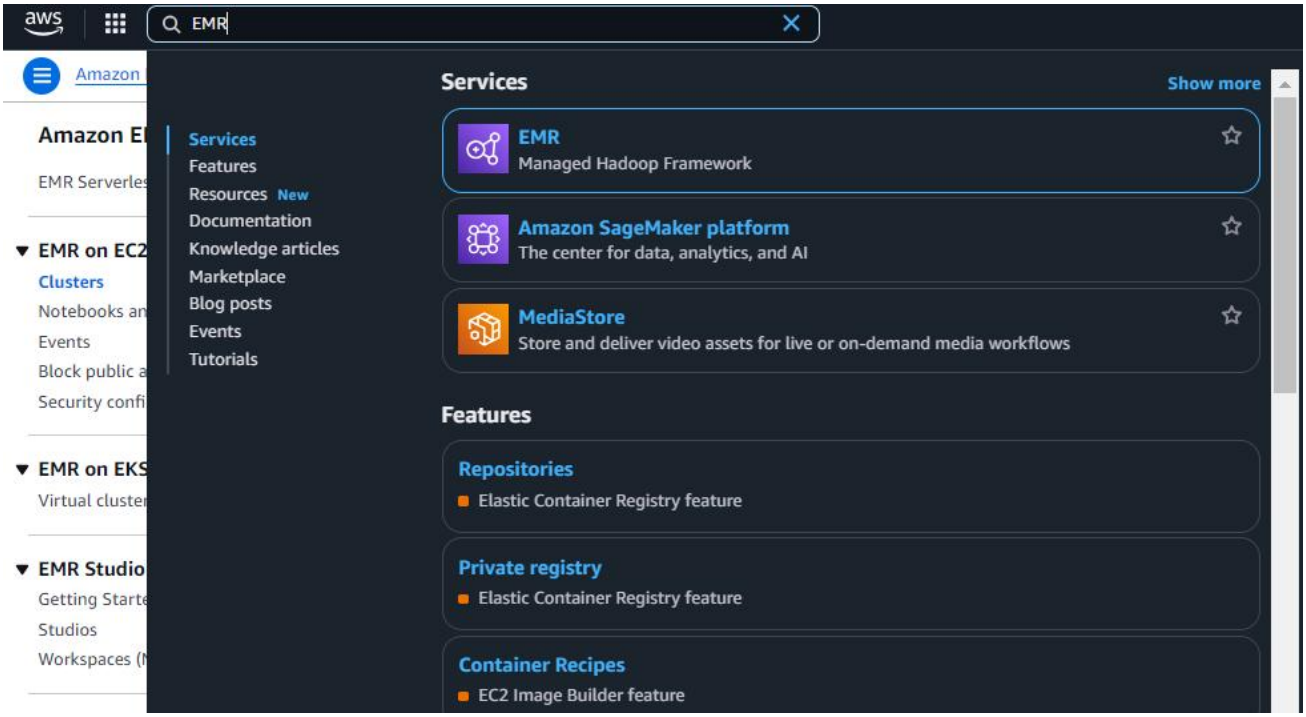

Follow the below screen shots to sucessfully create the SPARK Cluster

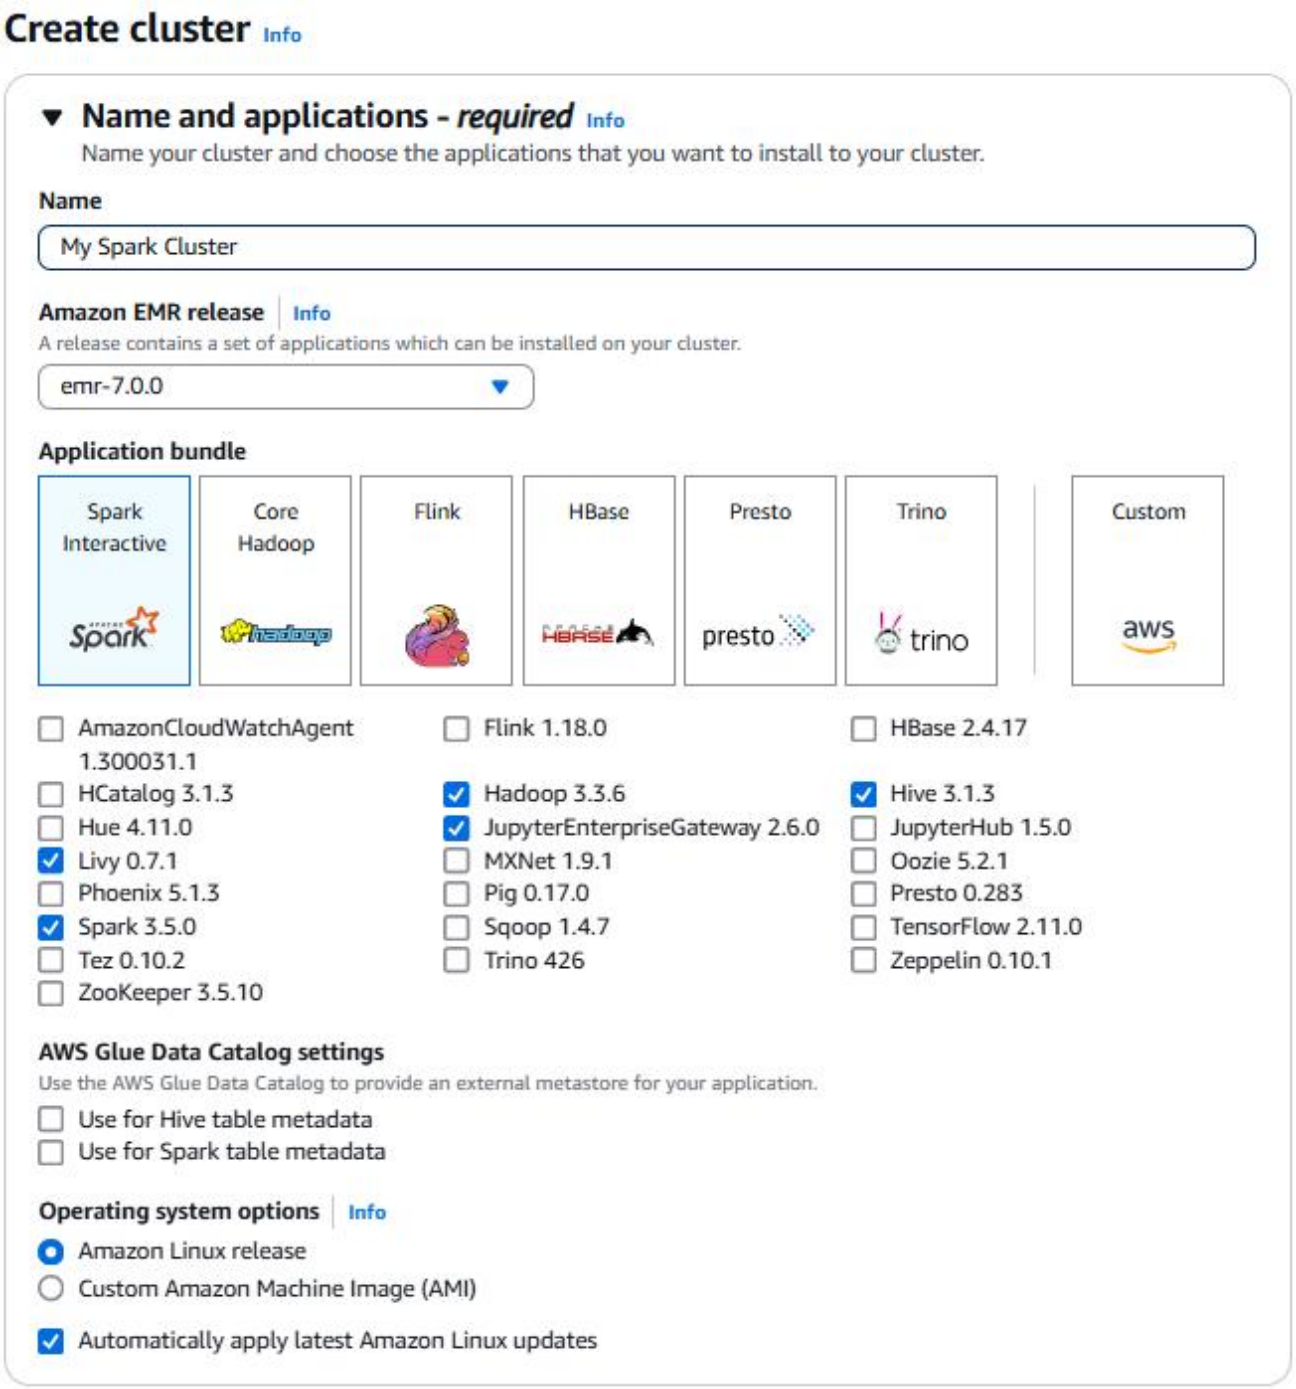

## Cluster configuration - *required* Info

Choose a configuration method for the primary, core, and task node groups for your cluster.

- (●) **Uniform instance groups**
  Choose the same EC2 instance type and purchasing option (On-Demand or Spot) for all nodes in your node group. Learn more
- ( ) **Flexible instance fleets**
  Choose from the widest variety of provisioning options for the EC2 instances in your cluster. Diversify instance types and purchasing options, and use an allocation strategy. Learn more

### Uniform instance groups

#### Primary

**Choose EC2 instance type**

m5.xlarge
4 vCore 16 GiB memory
EBS only storage On-Demand price: -
Lowest Spot price: -

Actions

☐ **Use high availability**
Launch highly available, more resilient cluster with three primary nodes on On-Demand Instances. This configuration applies for the lifetime of your cluster. Learn more

▶ **Node configuration - *optional***

#### Core

**Choose EC2 instance type**

m5.xlarge
4 vCore 16 GiB memory
EBS only storage On-Demand price: -
Lowest Spot price: -

Actions

▶ **Node configuration - *optional***

#### Task 1 of 1

Remove instance group

**Name**

Task - 1

**Choose EC2 instance type**

m5.xlarge
4 vCore 16 GiB memory
EBS only storage On-Demand price: -
Lowest Spot price: -

Actions

▶ **Node configuration - *optional***

### EBS root volume

EBS root volume applies to the operating systems and applications that you install on the cluster. EBS root volume ratio constraints

| Size (GiB) | IOPS | Throughput (MiB/s) |
|---|---|---|
| 15 | 3000 | 125 |
| 15 - 100 GiB per volume General Purpose SSD (gp3) | 3000 - 16000 IOPS per volume. Choose a maximum ratio of 500:1 between IOPS and volume size. | 125 - 1000 MiB/s per volume. Choose a maximum ratio of 0.25:1 between throughput and IOPS. |

## ▼ Cluster scaling and provisioning - *required* Info

Choose how Amazon EMR should size your cluster.

Choose an option

- (●) **Set cluster size manually** — Use this option if you know your workload patterns in advance.
- ( ) **Use EMR-managed scaling** — Monitor key workload metrics so that EMR can optimize the cluster size and resource utilization.
- ( ) **Use custom automatic scaling** — To programmatically scale core and task nodes, create custom automatic scaling policies.

### Provisioning configuration

Set the size of your core and task instance groups. Amazon EMR attempts to provision this capacity when you launch your cluster.

| Name | Instance type | Instance(s) size | Use Spot purchasing option |
|---|---|---|---|
| Core | m5.xlarge | 1 | ☐ |
| Task - 1 | m5.xlarge | 3 | ☐ |

## ▼ Networking - *required* Info

Choose the network settings that determine how you and other entities communicate with your cluster.

**Virtual private cloud (VPC)** Info

vpc-000e731e716efbe8e [Browse] [Create VPC]

**Subnet** Info

subnet-003cfbb9b5b0134d2 [Browse] [Create subnet]

▶ **EC2 security groups (firewall)**

## Cluster termination and node replacement Info

Choose termination settings and protect your cluster from accidental shutdown.

**Termination option**

- (•) Manually terminate cluster
- ( ) Automatically terminate cluster after last step ends
- ( ) Automatically terminate cluster after idle time (Recommended)

[x] Use termination protection

Protects your cluster from accidental termination. If on, you must first turn off protection to terminate the cluster. We recommend turning on termination protection for your long running clusters.

ⓘ To ensure unhealthy node replacement doesn't affect your existing workflows on EMR releases 7.0.0 and lower, we turn it off when you enable termination protection. You can change this setting when creating a cluster or by going to cluster configuration.

Unhealthy node replacement - *new* | Info

- ( ) Turn on

Amazon EMR gracefully stops processes on unhealthy nodes to minimize data loss and job interruptions. It quickly replaces unhealthy nodes with new EC2 instances to keep your jobs running smoothly.

- (•) Turn off

Amazon EMR adds unhealthy nodes to a denylist while keeping them in the cluster, allowing you continued access for troubleshooting.

## Bootstrap actions (0) Info

Remove | Edit | Add

Use bootstrap actions to install software or customize your instance configuration.

## Cluster logs Info

Choose where and how to store your log files.

## Tags Info

Use tags to search and filter for resources, and track AWS costs associated with your cluster.

## Software settings Info

Override the default configurations for specific applications on your cluster.

## Security configuration and EC2 key pair Info

Choose a security configuration or create a new one that you can reuse with other clusters.

**Security configuration**

Select your cluster encryption, authentication, authorization, and instance metadata service settings.

Choose a security configuration | Browse | Create security configuration

**Amazon EC2 key pair for SSH to the cluster** Info

vockey | Browse | Create key pair

## Identity and Access Management (IAM) roles - *required* Info

Choose or create a service role and instance profile for the EC2 instances in your cluster.

### Amazon EMR service role Info

The service role is an IAM role that Amazon EMR assumes to provision resources and perform service-level actions with other AWS services.

- (●) Choose an existing service role
  Select a default service role or a custom role with IAM policies attached so that your cluster can interact with other AWS services.
- ( ) Create a service role
  Let Amazon EMR create a new service role so that you can grant and restrict access to resources in other AWS services.

**Service role**

EMR_DefaultRole

### EC2 instance profile for Amazon EMR

The instance profile assigns a role to every EC2 instance in a cluster. The instance profile must specify a role that can access the resources for your steps and bootstrap actions.

- (●) Choose an existing instance profile
  Select a default role or a custom instance profile with IAM policies attached so that your cluster can interact with your resources in Amazon S3.
- ( ) Create an instance profile
  Let Amazon EMR create a new instance profile so that you can specify a custom set of resources for it to access in Amazon S3.

**Instance profile**

EMR_EC2_DefaultRole

### Custom automatic scaling role - *optional*

When a custom automatic scaling rule triggers, Amazon EMR assumes this role to add and terminate EC2 instances. Learn more

**Custom automatic scaling role**

EMR_AutoScaling_DefaultRole | Create IAM role

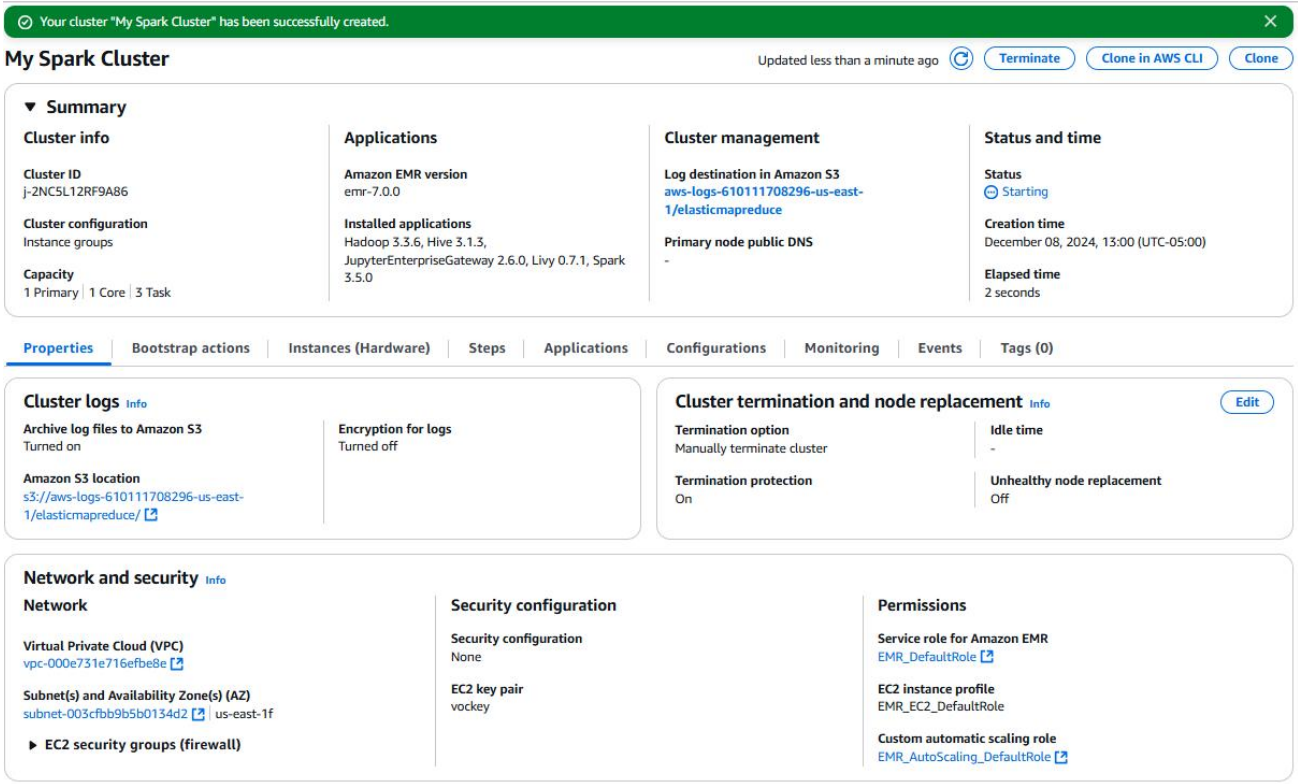
Your cluster "My Spark Cluster" has been successfully created.
My Spark Cluster
Updated less than a minute ago
Terminate
Clone in AWS CLI
Clone
Summary
Cluster info
Cluster ID
j-2NC5L12RF9A86
Cluster configuration
Instance groups
Capacity
1 Primary | 1 Core | 3 Task
Applications
Amazon EMR version
emr-7.0.0
Installed applications
Hadoop 3.3.6, Hive 3.1.3, JupyterEnterpriseGateway 2.6.0, Livy 0.7.1, Spark 3.5.0
Cluster management
Log destination in Amazon S3
aws-logs-610111708296-us-east-1/elasticmapreduce
Primary node public DNS
-
Status and time
Status
Starting
Creation time
December 08, 2024, 13:00 (UTC-05:00)
Elapsed time
2 seconds
Properties
Bootstrap actions
Instances (Hardware)
Steps
Applications
Configurations
Monitoring
Events
Tags (0)
Cluster logs Info
Archive log files to Amazon S3
Turned on
Amazon S3 location
s3://aws-logs-610111708296-us-east-1/elasticmapreduce/
Encryption for logs
Turned off
Cluster termination and node replacement Info
Edit
Termination option
Manually terminate cluster
Termination protection
On
Idle time
-
Unhealthy node replacement
Off
Network and security Info
Network
Virtual Private Cloud (VPC)
vpc-000e731e716efbe8e
Subnet(s) and Availability Zone(s) (AZ)
subnet-003cfbb9b5b0134d2 | us-east-1f
EC2 security groups (firewall)
Security configuration
Security configuration
None
EC2 key pair
vockey
Permissions
Service role for Amazon EMR
EMR_DefaultRole
EC2 instance profile
EMR_EC2_DefaultRole
Custom automatic scaling role
EMR_AutoScaling_DefaultRole

## Step-3: Modify the security rules for the cluster

Now needs to update the inbound security rules

Do to the Network and Security -> Expand EC2 security groups -> click on core and task nodes

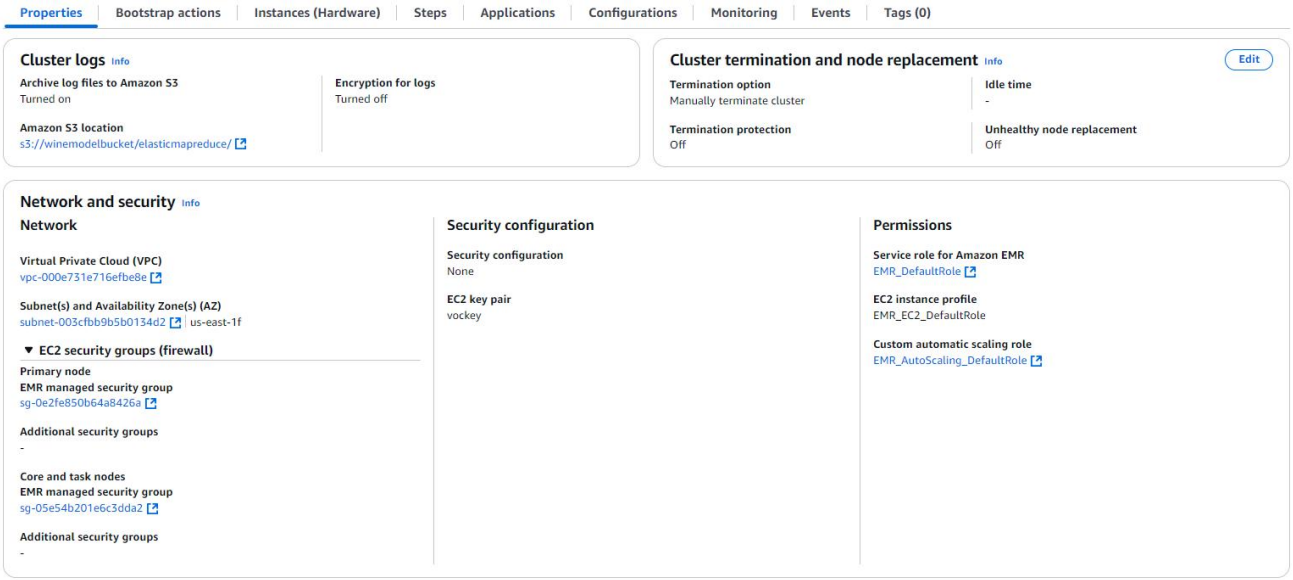


Click on edit inbound rules and add as shown in screenshot

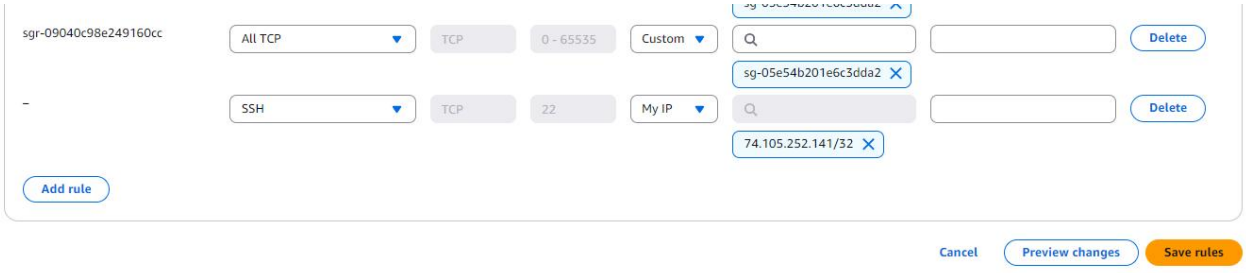

## Step-4: Login to the master node and transfer all the required files into it

After selecting master node click on connect and

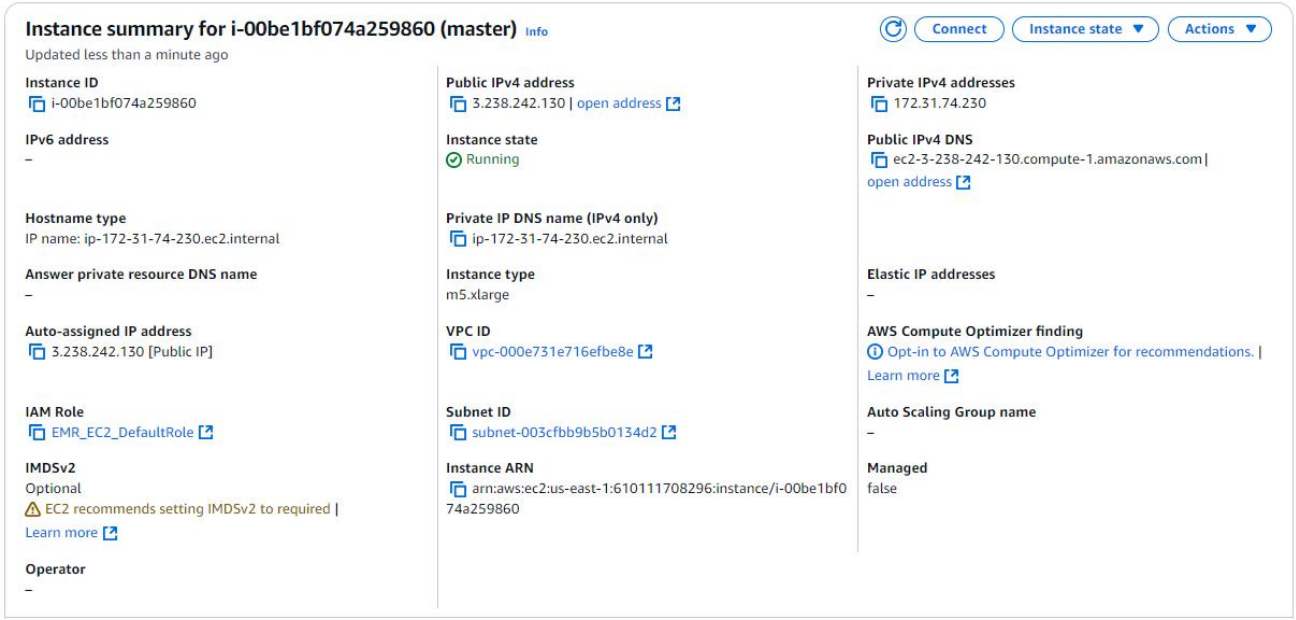


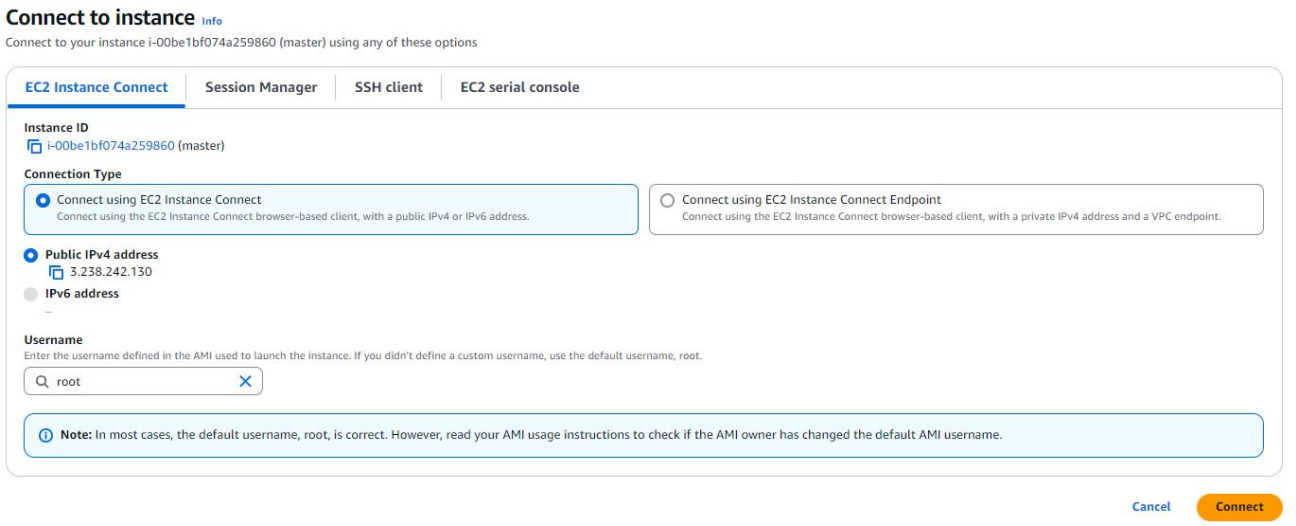


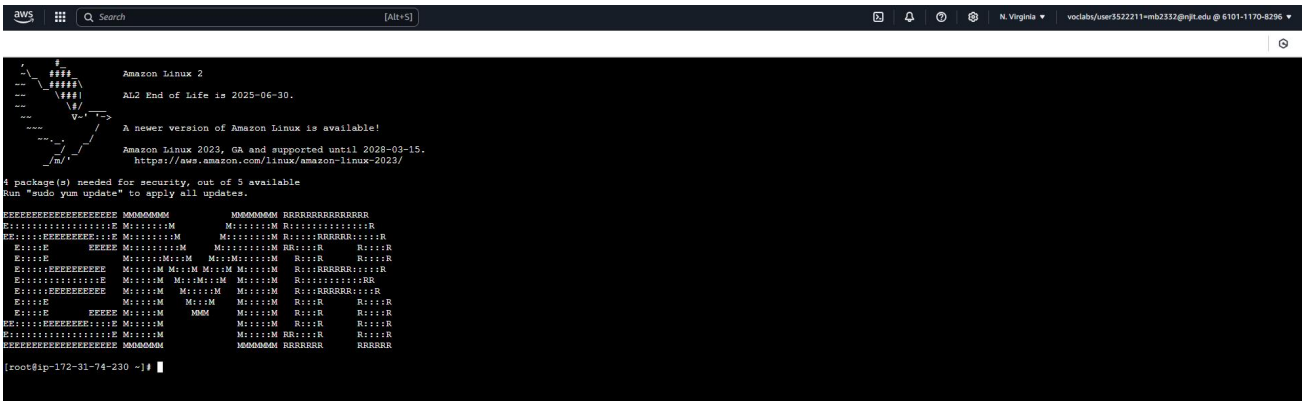

## Step-5: Execute the training.py ML application

Enter command to install git and enter yes when ask
sudo yum install git

Now clone with the github repository
git clone git clone https://github.com/Manas1227/cs643-853-pa2-mb2332

Now change to the new directory cs643-853-pa2-mb2332 and check it clone correctly
cd cs643-853-pa2-mb2332
ls

```
[root@ip-172-31-74-230 ~]# git clone https://github.com/Manas1227/cs643-853-pa2-mb2332
Cloning into 'cs643-853-pa2-mb2332'...
Username for 'https://github.com': Manas1227
Password for 'https://Manas1227@github.com':
remote: Enumerating objects: 39, done.
remote: Counting objects: 100% (39/39), done.
remote: Compressing objects: 100% (28/28), done.
remote: Total 39 (delta 15), reused 27 (delta 8), pack-reused 0 (from 0)
Receiving objects: 100% (39/39), 35.23 KiB | 17.62 MiB/s, done.
Resolving deltas: 100% (15/15), done.
[root@ip-172-31-74-230 ~]# ls
cs643-853-pa2-mb2332
[root@ip-172-31-74-230 ~]# cd cs643-853-pa2-mb2332/
[root@ip-172-31-74-230 cs643-853-pa2-mb2332]# ls
LICENSE  prediction.py  README.md  requirements.txt  TrainingDataset.csv  training.py  ValidationDataset.csv
[root@ip-172-31-74-230 cs643-853-pa2-mb2332]#
```

Now install all required libraries by followed command
sudo pip3 install -r requirements.txt

Run the training file
sudo spark-submit training.py

```
24/12/11 22:48:17 INFO FileFormatWriter: Finished processing stats for write job f765c4db-03e2-4e2d-9215-eb36cf8cfb35.
24/12/11 22:48:17 INFO Instrumentation: [ce5f8861] training finished
24/12/11 22:48:17 INFO Instrumentation: [be1614d4] training finished
Training completed in 64.95 seconds
{'Model': 'RandomForestClassifier', 'Accuracy': 0.95625, 'Recall': 0.95625, 'F1 Score': 0.9447916666666667}
{'Model': 'LogisticRegression', 'Accuracy': 0.975, 'Recall': 0.9750000000000001, 'F1 Score': 0.9729166666666667}
{'Model': 'DecisionTreeClassifier', 'Accuracy': 1.0, 'Recall': 0.9999999999999999, 'F1 Score': 0.9999999999999999}
24/12/11 22:48:17 INFO SparkContext: SparkContext is stopping with exitCode 0.
24/12/11 22:48:17 INFO SparkUI: Stopped Spark web UI at http://ip-172-31-69-53.ec2.internal:4040
```

Training model will train the ML model using EMR cluster and upload the best model to S3 bucket "trainingmodel.model"

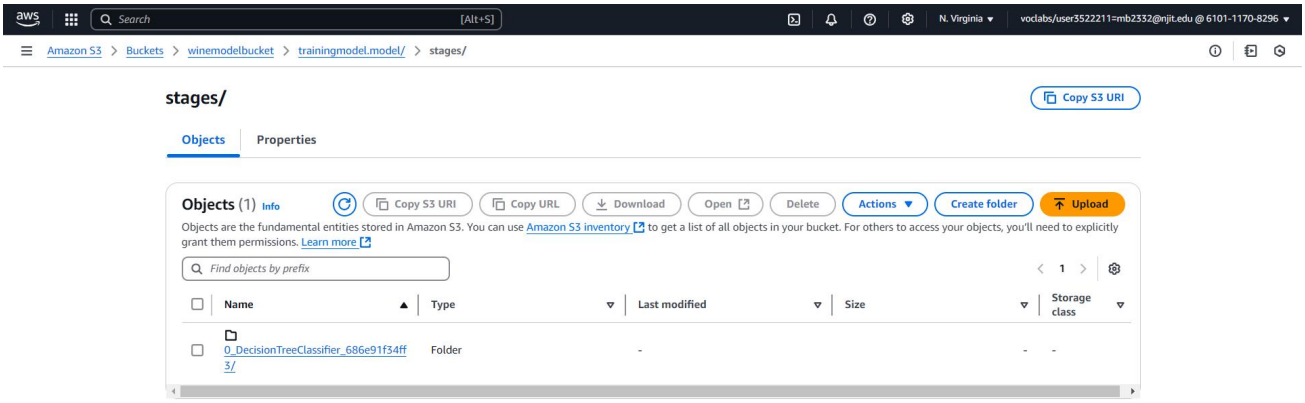

**Step-6: Terminate the EMR cluster**

Terminate the EMR cluster after sucessfully competed our training part
If asked **Turn off** the protection and **Terminate** the cluster

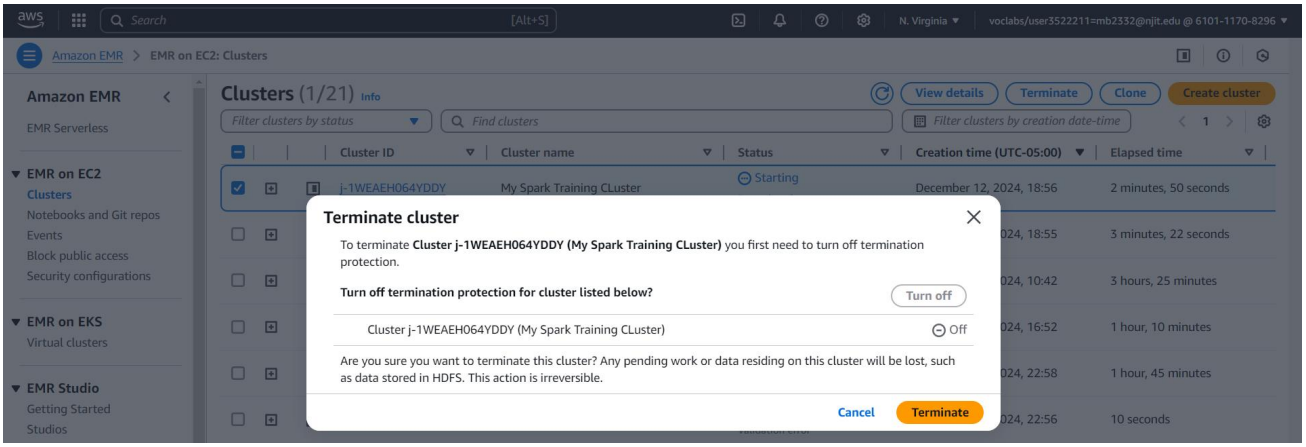

# SECTION 2: AWS Cloud setup for running the prediction ML application - without Docker

Step-1: Create standalone EC2 instance on AWS

Redirect to **EC2** on AWS console and click on **Launch Instance**

Enter Name and Select Amazon Linux-2023 AMI

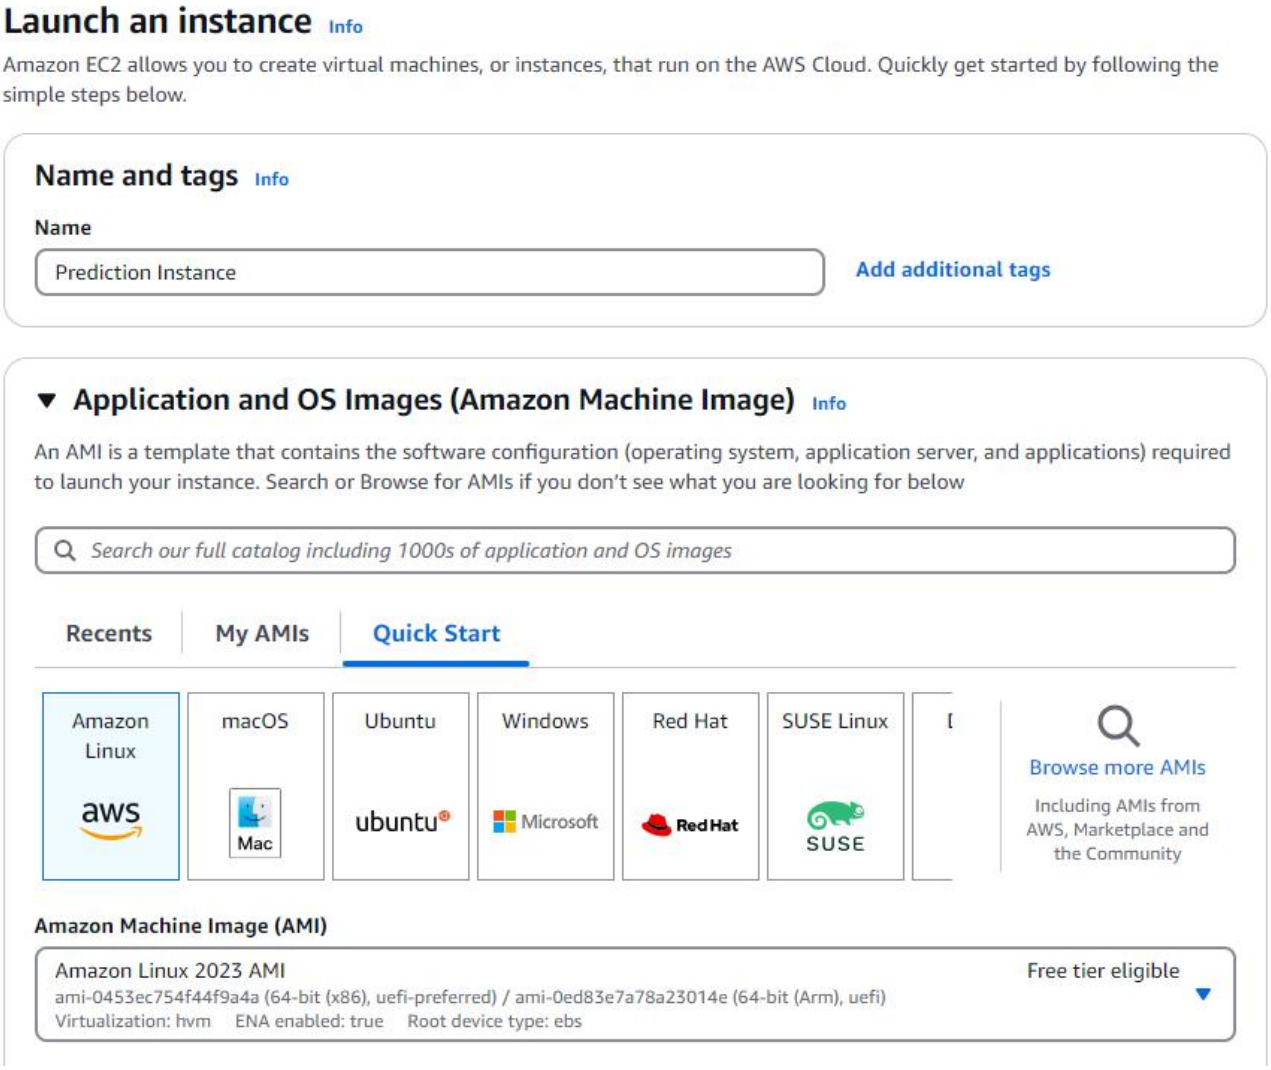

Select t2.medium as a Instance type

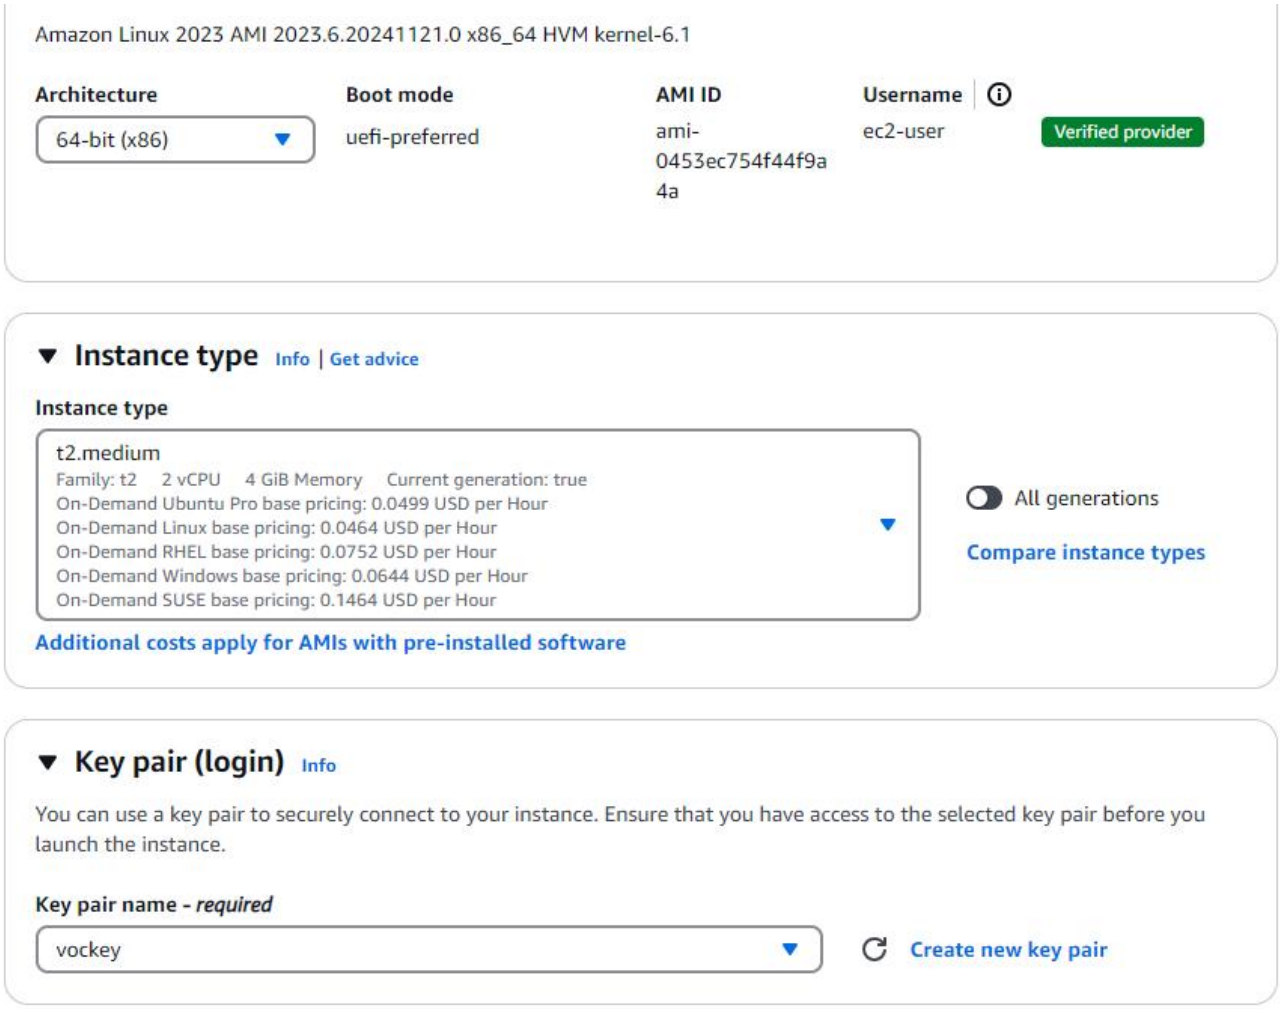

Create or choose from existing security group to inbound rule to accept ssh connection from MyIp
Also configure with 15 GB storage as we are going to install all the required dependencies

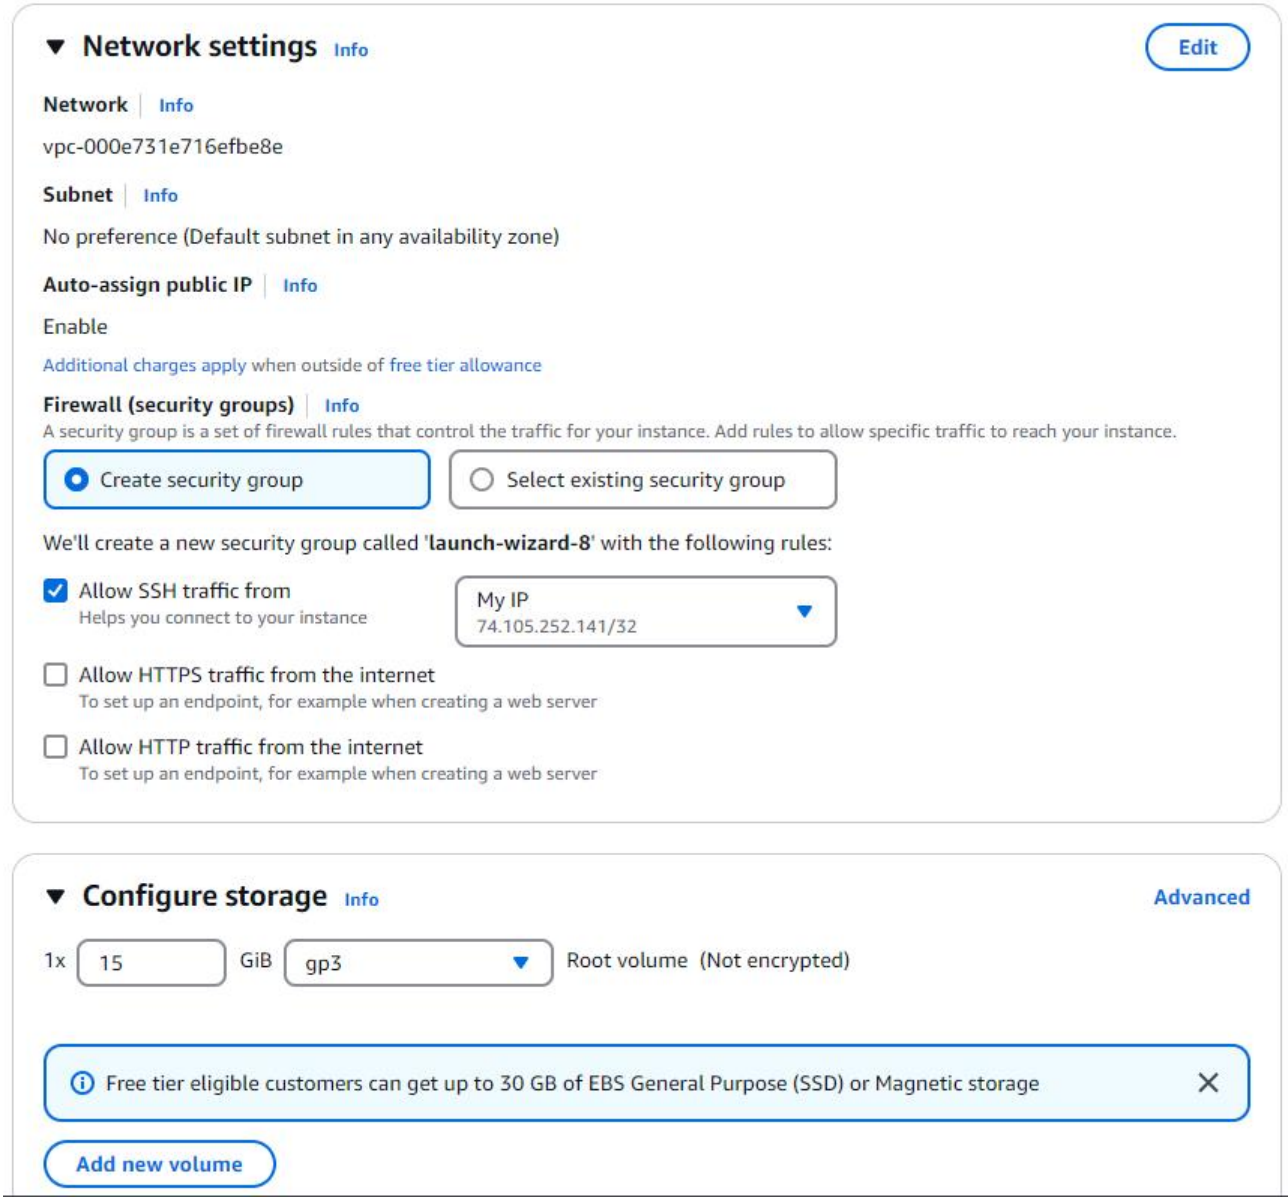





Make sure to select IAM role with all required AWS services access

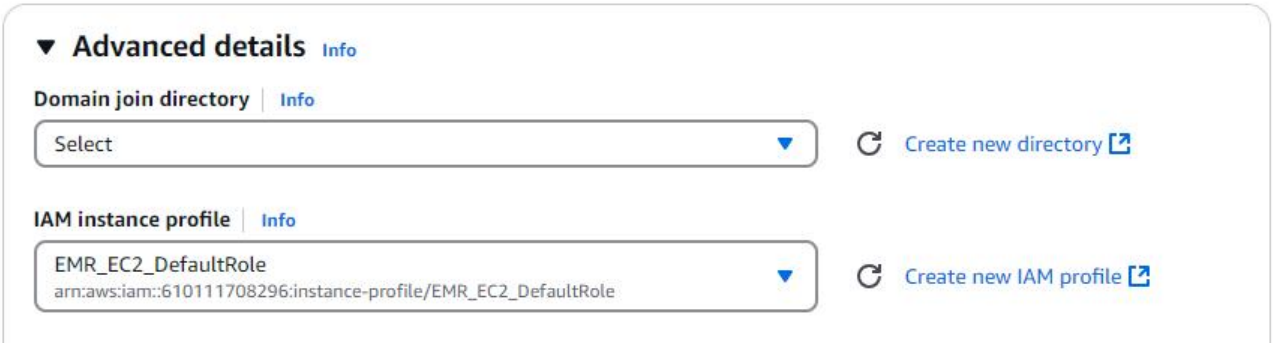

Step-2: Login to the EC2 instance and setup environment

Click on Connect and connect to the instance with your preferred options

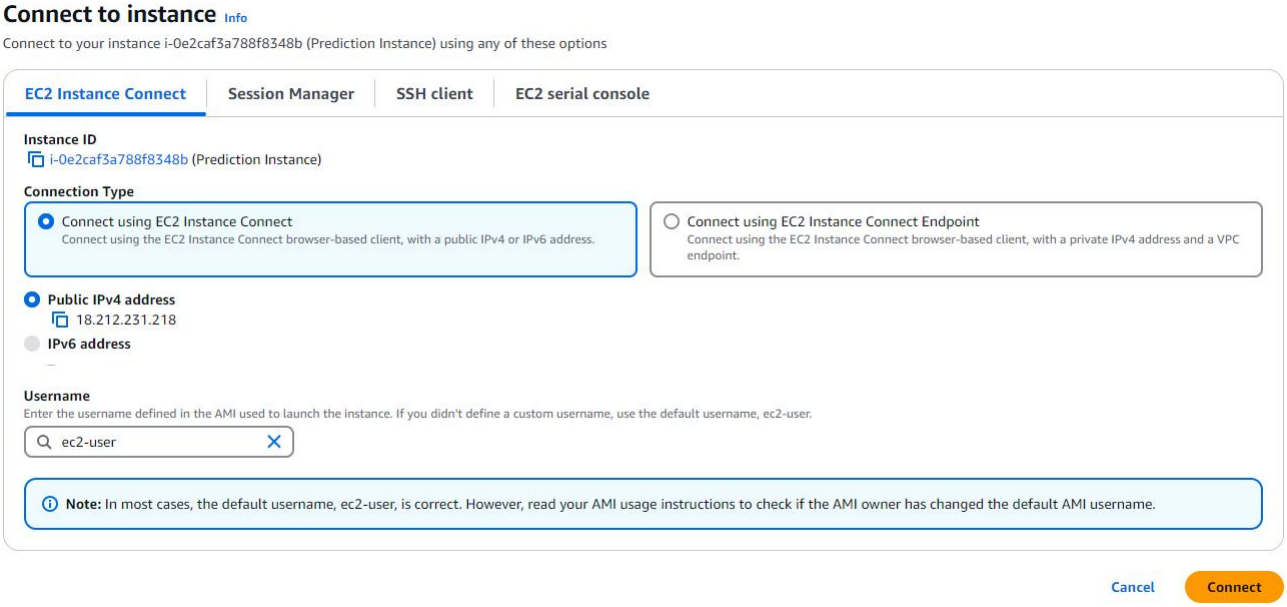


Now install Github to clone with the repository
sudo yum install git

Clone with the github repository to access all the program files on EC2 instance
git clone https://github.com/Manas1227/cs643-853-pa2-mb2332

```
[ec2-user@ip-172-31-19-212 ~]$ git clone https://github.com/Manas1227/cs643-853-pa2-mb2332
Cloning into 'cs643-853-pa2-mb2332'...
Username for 'https://github.com': Manas1227
Password for 'https://Manas1227@github.com':
remote: Enumerating objects: 51, done.
remote: Counting objects: 100% (51/51), done.
remote: Compressing objects: 100% (38/38), done.
remote: Total 51 (delta 19), reused 36 (delta 9), pack-reused 0 (from 0)
Receiving objects: 100% (51/51), 37.72 KiB | 18.86 MiB/s, done.
Resolving deltas: 100% (19/19), done.
[ec2-user@ip-172-31-19-212 ~]$
```

Go to the new project directory
cd cs643-853-pa2-mb2332/

Check the Python and pip version and if not install them on the EC2
python3 --version
python3.9 -m ensurepip --upgrade

```
[ec2-user@ip-172-31-19-212 cs643-853-pa2-mb2332]$ python3 --version
Python 3.9.16
[ec2-user@ip-172-31-19-212 cs643-853-pa2-mb2332]$ python3.9 -m ensurepip --upgrade
Defaulting to user installation because normal site-packages is not writeable
Looking in links: /tmp/tmpls0wz7o1
Requirement already satisfied: setuptools in /usr/lib/python3.9/site-packages (59.6.0)
Processing /tmp/tmpls0wz7o1/pip-21.3.1-py3-none-any.whl
Installing collected packages: pip
Successfully installed pip-21.3.1
[ec2-user@ip-172-31-19-212 cs643-853-pa2-mb2332]$
```

Install all the requirements listed in requirements.txt file

pip3 install -r requirements.txt

```
[ec2-user@ip-172-31-19-212 cs643-853-pa2-mb2332]$ pip3 install -r requirements.txt
Defaulting to user installation because normal site-packages is not writeable
Collecting pyspark
  Downloading pyspark-3.5.3.tar.gz (317.3 MB)
     |                                | 286.9 MB 10 kB/s eta 0:46:26
     |                                | 317.3 MB 38 kB/s
  Preparing metadata (setup.py) ... done
Collecting numpy
  Downloading numpy-2.0.2-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (19.5 MB)
     |                                | 19.5 MB 10 kB/s
Collecting findspark
  Downloading findspark-2.0.1-py2.py3-none-any.whl (4.4 kB)
Collecting py4j==0.10.9.7
  Downloading py4j-0.10.9.7-py2.py3-none-any.whl (200 kB)
     |                                | 200 kB 93.5 MB/s
Using legacy 'setup.py install' for pyspark, since package 'wheel' is not installed.
Installing collected packages: py4j, pyspark, numpy, findspark
    Running setup.py install for pyspark ... done
Successfully installed findspark-2.0.1 numpy-2.0.2 py4j-0.10.9.7 pyspark-3.5.3
WARNING: You are using pip version 21.3.1; however, version 24.3.1 is available.
You should consider upgrading via the '/usr/bin/python3.9 -m pip install --upgrade pip' command.
[ec2-user@ip-172-31-19-212 cs643-853-pa2-mb2332]$
```

Install java on EC2 instance

sudo yum install java-11-amazon-corretto-devel

Add Java home and path variables in below file and save it

nano ~/.bashrc

export JAVA_HOME=/usr/lib/jvm/java-11-amazon-corretto.x86_64

export PATH=$PATH:$JAVA_HOME/bin

Now download required jars

wget https://repo1.maven.org/maven2/org/apache/hadoop/hadoop-aws/3.3.1/hadoop-aws-3.3.1.jar -P ~/libs/

wget https://repo1.maven.org/maven2/com/amazonaws/aws-java-sdk-bundle/1.11.1000/aws-java-sdk-bundle-1.11.1000.jar -P ~/libs/

```
[ec2-user@ip-172-31-19-212 ~]$ wget https://repo1.maven.org/maven2/org/apache/hadoop/hadoop-aws/3.3.1/hadoop-aws-3.3.1.jar -P ~/libs/
--2024-12-13 01:16:38--  https://repo1.maven.org/maven2/org/apache/hadoop/hadoop-aws/3.3.1/hadoop-aws-3.3.1.jar
Resolving repo1.maven.org (repo1.maven.org)... 146.75.32.209, 2a04:4e42:79::209
Connecting to repo1.maven.org (repo1.maven.org)|146.75.32.209|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 870644 (850K) [application/java-archive]
Saving to: '/home/ec2-user/libs/hadoop-aws-3.3.1.jar'

hadoop-aws-3.3.1.jar                100%[=========================================================================================>] 850.24K  --.-KB/s    in 0.02s

2024-12-13 01:16:38 (42.3 MB/s) - '/home/ec2-user/libs/hadoop-aws-3.3.1.jar' saved [870644/870644]

[ec2-user@ip-172-31-19-212 ~]$ wget https://repo1.maven.org/maven2/com/amazonaws/aws-java-sdk-bundle/1.11.1000/aws-java-sdk-bundle-1.11.1000.jar -P ~/libs/
--2024-12-13 01:17:33--  https://repo1.maven.org/maven2/com/amazonaws/aws-java-sdk-bundle/1.11.1000/aws-java-sdk-bundle-1.11.1000.jar
Resolving repo1.maven.org (repo1.maven.org)... 146.75.32.209, 2a04:4e42:79::209
Connecting to repo1.maven.org (repo1.maven.org)|146.75.32.209|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 220787798 (211M) [application/java-archive]
Saving to: '/home/ec2-user/libs/aws-java-sdk-bundle-1.11.1000.jar'

aws-java-sdk-bundle-1.11.1000.jar   100%[=========================================================================================>] 210.56M  72.1MB/s    in 2.9s

2024-12-13 01:17:36 (72.1 MB/s) - '/home/ec2-user/libs/aws-java-sdk-bundle-1.11.1000.jar' saved [220787798/220787798]

[ec2-user@ip-172-31-19-212 ~]$
```

Now run the prediction.py file

spark-submit --jars /home/ec2-user/libs/hadoop-aws-3.3.1.jar,/home/ec2-user/libs/aws-java-sdk-bundle-1.11.1000.jar prediction.py

```
24/12/13 01:19:48 INFO DAGScheduler: Job 11 is finished. Cancelling potential speculative or zombie tasks for this job
24/12/13 01:19:48 INFO TaskSchedulerImpl: Killing all running tasks in stage 14: Stage finished
24/12/13 01:19:48 INFO DAGScheduler: Job 11 finished: collectAsMap at MulticlassMetrics.scala:61, took 0.260521 s
Test Accuracy: 1.0
Test F1 Score: 0.9999999999999999
24/12/13 01:19:48 INFO SparkContext: SparkContext is stopping with exitCode 0.
24/12/13 01:19:48 INFO BlockManagerInfo: Removed broadcast_18_piece0 on ip-172-31-19-212.ec2.internal:41931 in memory (size: 34.8 KiB, free: 434.2 MiB)
24/12/13 01:19:48 INFO SparkUI: Stopped Spark web UI at http://ip-172-31-19-212.ec2.internal:4040
24/12/13 01:19:48 INFO BlockManagerInfo: Removed broadcast_19_piece0 on ip-172-31-19-212.ec2.internal:41931 in memory (size: 35.9 KiB, free: 434.3 MiB)
24/12/13 01:19:48 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
24/12/13 01:19:48 INFO MemoryStore: MemoryStore cleared
24/12/13 01:19:48 INFO BlockManager: BlockManager stopped
24/12/13 01:19:48 INFO BlockManagerMaster: BlockManagerMaster stopped
24/12/13 01:19:48 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
24/12/13 01:19:48 INFO SparkContext: Successfully stopped SparkContext
24/12/13 01:19:48 INFO ShutdownHookManager: Shutdown hook called
24/12/13 01:19:48 INFO ShutdownHookManager: Deleting directory /tmp/spark-0507b134-be05-4adc-b82d-15a624e8e2cf/pyspark-ef268633-6e48-44d9-adfe-21850440a779
24/12/13 01:19:48 INFO ShutdownHookManager: Deleting directory /tmp/spark-aae559ca-56bf-47d8-a10d-250ff4a2082a
24/12/13 01:19:48 INFO ShutdownHookManager: Deleting directory /tmp/spark-0507b134-be05-4adc-b82d-15a624e8e2cf
24/12/13 01:19:48 INFO MetricsSystemImpl: Stopping s3a-file-system metrics system...
24/12/13 01:19:48 INFO MetricsSystemImpl: s3a-file-system metrics system stopped.
24/12/13 01:19:48 INFO MetricsSystemImpl: s3a-file-system metrics system shutdown complete.
```

As shown above, DecisionTree Model from s3 was used to predict the label for ValidationDataset.csv

# SECTION 3: Create Docker image and push it on Docker repository

**Step-1: Create docker image by following command:**

docker build -t aws-spark-training .

docker images

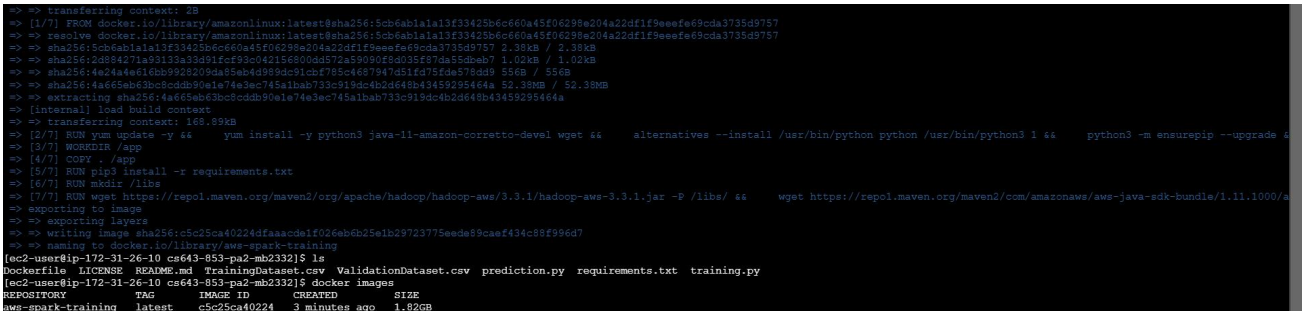


**Step-2: Upload docker image to docker repository**

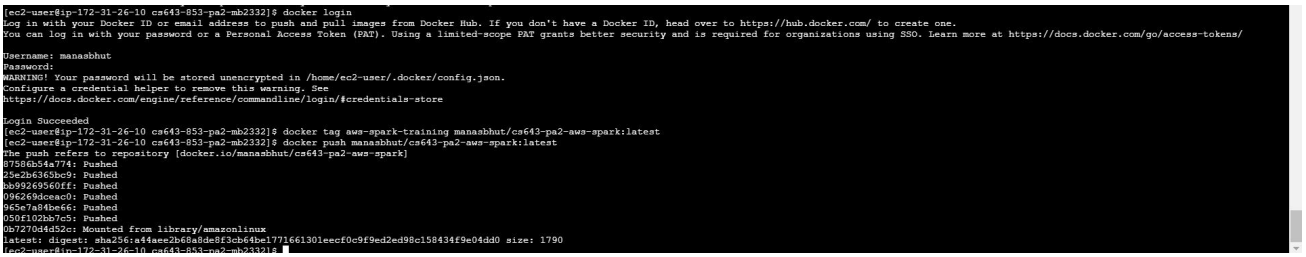


**Step-3: Go to docker repository to verify the uploaded image**

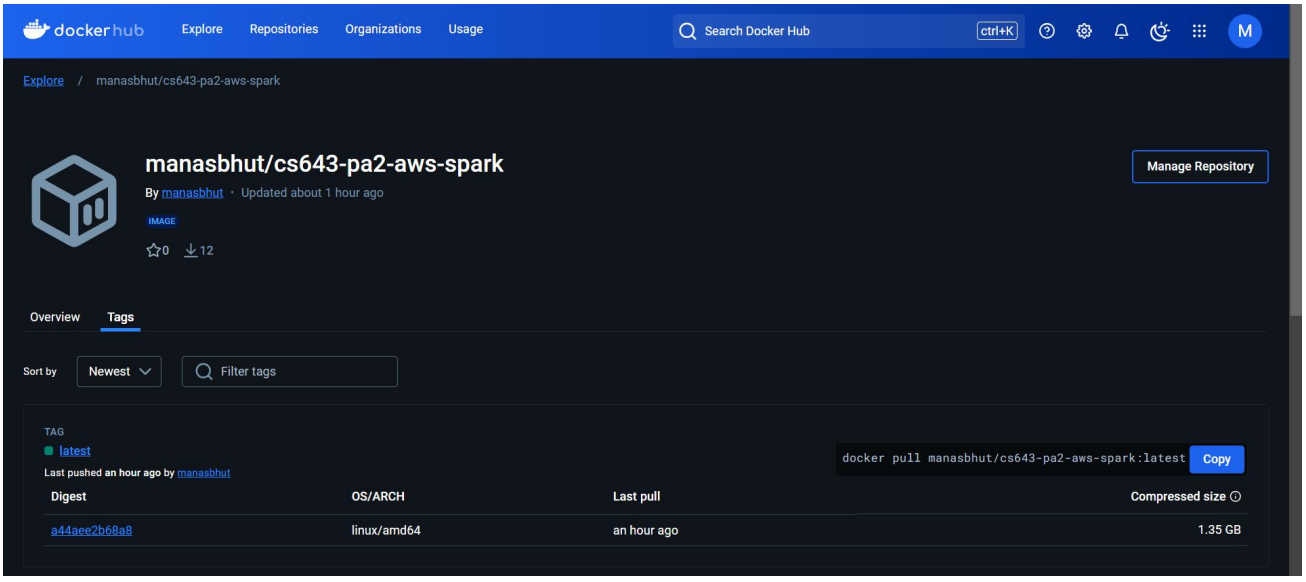

## SECTION 4: Run docker image on newly created EC2 instance

**Step-1: Install and run docker on EC2 instance**

sudo yum install docker -y
sudo service docker start
sudo usermod -aG docker ec2-user

**Step-2: Verify by running below command**

docker info

```
   ,     #_
   ~\_  ####_        Amazon Linux 2023
  ~~  \_#####\
  ~~     \###|
  ~~       \#/ ___   https://aws.amazon.com/linux/amazon-linux-2023
   ~~       V~' '->
    ~~~         /
      ~~._.   _/
         _/ _/
       _/m/'
Last login: Thu Dec 12 23:19:02 2024 from 18.206.107.27
[ec2-user@ip-172-31-24-30 ~]$ docker info
Client:
 Version:    25.0.5
 Context:    default
 Debug Mode: false
 Plugins:
  buildx: Docker Buildx (Docker Inc.)
    Version:  v0.0.0+unknown
    Path:     /usr/libexec/docker/cli-plugins/docker-buildx

Server:
 Containers: 0
  Running: 0
  Paused: 0
  Stopped: 0
 Images: 0
 Server Version: 25.0.6
 Storage Driver: overlay2
  Backing Filesystem: xfs
  Supports d_type: true
  Using metacopy: false
  Native Overlay Diff: true
  userxattr: false
 Logging Driver: json-file
 Cgroup Driver: systemd
 Cgroup Version: 2
 Plugins:
  Volume: local
  Network: bridge host ipvlan macvlan null overlay
  Log: awslogs fluentd gcplogs gelf journald json-file local splunk syslog
 Swarm: inactive
 Runtimes: io.containerd.runc.v2 runc
 Default Runtime: runc
```

If it still giving error than try to exit the Ec2 instance and login again

exit

Now login again to the Ec2 instance and verify docker is up and running

docker info

Pull the latest image from Docker repository
docker pull manasbhut/cs643-pa2-aws-spark:latest
docker images

```
[ec2-user@ip-172-31-24-30 ~]$ docker pull manasbhut/cs643-pa2-aws-spark:latest
latest: Pulling from manasbhut/cs643-pa2-aws-spark
Digest: sha256:a44aee2b68a8de8f3cb64be1771661301eecf0c9f9ed2ed98c158434f9e04dd0
Status: Image is up to date for manasbhut/cs643-pa2-aws-spark:latest
docker.io/manasbhut/cs643-pa2-aws-spark:latest
[ec2-user@ip-172-31-24-30 ~]$ docker images
REPOSITORY                          TAG       IMAGE ID       CREATED        SIZE
manasbhut/cs643-pa2-aws-spark       latest    c5c25ca40224   3 hours ago    1.82GB
```

Run the pulled image by following command
docker run -d --name spark-container manasbhut/cs643-pa2-aws-spark:latest
docker ps -a

```
[ec2-user@ip-172-31-24-30 ~]$ docker ps -a
CONTAINER ID   IMAGE                                COMMAND                  CREATED          STATUS                      PORTS     NAMES
fc3a7b4073d4   manasbhut/cs643-pa2-aws-spark:latest   "spark-submit --jars…"   49 seconds ago   Exited (0) 13 seconds ago             spark-container
[ec2-user@ip-172-31-24-30 ~]$
```

Run following command to see the console output
docker logs spark-container

```
24/12/13 01:29:48 INFO DAGScheduler: ResultStage 14 (collectAsMap at MulticlassMetrics.scala:61) finished in 0.029 s
24/12/13 01:29:48 INFO DAGScheduler: Job 11 is finished. Cancelling potential speculative or zombie tasks for this job
24/12/13 01:29:48 INFO TaskSchedulerImpl: Removed TaskSet 14.0, whose tasks have all completed, from pool
24/12/13 01:29:48 INFO TaskSchedulerImpl: Killing all running tasks in stage 14: Stage finished
24/12/13 01:29:48 INFO DAGScheduler: Job 11 finished: collectAsMap at MulticlassMetrics.scala:61, took 0.241529 s
Test Accuracy: 1.0
Test F1 Score: 0.9999999999999999
24/12/13 01:29:48 INFO SparkContext: SparkContext is stopping with exitCode 0.
24/12/13 01:29:48 INFO SparkUI: Stopped Spark web UI at http://fc3a7b4073d4:4040
24/12/13 01:29:48 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
24/12/13 01:29:48 INFO MemoryStore: MemoryStore cleared
24/12/13 01:29:48 INFO BlockManager: BlockManager stopped
24/12/13 01:29:48 INFO BlockManagerMaster: BlockManagerMaster stopped
24/12/13 01:29:48 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
24/12/13 01:29:48 INFO SparkContext: Successfully stopped SparkContext
24/12/13 01:29:48 INFO ShutdownHookManager: Shutdown hook called
24/12/13 01:29:48 INFO ShutdownHookManager: Deleting directory /tmp/spark-2994c3de-7dd9-468d-8451-6f5040012a38
24/12/13 01:29:48 INFO ShutdownHookManager: Deleting directory /tmp/spark-d5d94d9f-6256-4475-8223-5ec1e1aa4c52/pyspark-7ffebe8a-6ded-4103-b027-d84fa771ea9c
24/12/13 01:29:48 INFO ShutdownHookManager: Deleting directory /tmp/spark-d5d94d9f-6256-4475-8223-5ec1e1aa4c52
24/12/13 01:29:48 INFO MetricsSystemImpl: Stopping s3a-file-system metrics system...
24/12/13 01:29:48 INFO MetricsSystemImpl: s3a-file-system metrics system stopped.
24/12/13 01:29:48 INFO MetricsSystemImpl: s3a-file-system metrics system shutdown complete.
[ec2-user@ip-172-31-24-30 ~]$ docker logs spark-container
```